

A Model of ChatGPT

Tom Cunningham

2025-06-14

Abstract

This note gives a simple formal model to predict when people will use ChatGPT. In brief: when someone encounters a question which has components that they do not have experience with, but which are represented in ChatGPT's training data (i.e., the public internet). This is based on a simple formal model where both humans and LLMs answer questions by interpolating among questions they already know the answer to. The model is an extension of an earlier [model of LLMs](#) I developed for a different purpose.

This note gives a simple formal model of ChatGPT use. ChatGPT is being used for an extraordinary variety of purposes, and this model clearly doesn't capture all of them, but I think it's a useful and precise baseline.

Suppose a human encounters a question, \mathbf{q} , which is a vector of binary characteristics. The true answer is a scalar, a , determined by a set of unobserved weights $a = \mathbf{w}\mathbf{q}$. The human guesses the answer to the new question by interpolating among previously-seen questions $((\mathbf{q}^i, a^i)_{i=1, \dots, n})$. They can also consult ChatGPT, which answers questions in the same way, but with a different set of previously-seen questions (i.e. the training data). We can then give a crisp closed-form expression for the expected value of consulting ChatGPT before answering a question (the expected "uplift"). The expected uplift will depend on a number of factors:

1. **The dimensionality of the question (p).** A higher-dimensional problem may be more costly to enter into ChatGPT, but it will also increase the potential benefit.
2. **Public experience.** These are the training questions that the AI has observed, which we can conceptualize as the corpus of public knowledge (e.g., the internet).
3. **Private experience.** These are the questions that the user has personally encountered and for which they have observed the true answer.

The model gives several predictions:

1. ChatGPT will not be used for questions the user has encountered before.
2. ChatGPT is more likely to be used for domains with higher *latent* dimensionality (p).
3. ChatGPT is more likely to be used for domains with lower *surface* dimensionality, as this reduces the cost of specifying the question.
4. ChatGPT is more likely to be used by humans with less experience in a domain (i.e., smaller n_{private}).

0.1 Conjectures About Adoption

We can make some conjectures about adoption by occupation and by task:

Occupation	Predicted ChatGPT use	Reason
software engineer - python	high	problems novel, discrete, similar to those on internet
software engineer - proprietary	low	problems novel, discrete, not similar to those on internet
physician	high	problems novel, discrete, similar to those on internet
contact center worker	low	problems novel, discrete, not similar to those on internet
architect	low	problems novel, not discrete, not text-based
manual worker	low	problems not text-based

Table 1: Conjectures about adoption by occupation

Task	Predicted ChatGPT use	Reason
Intellectual curiosity	high	novel, discrete, similar to those on the internet
Diagnosing medical problems	high	novel, discrete, similar to those on the internet
Problems with widely-adopted systems (car, house, computer)	high	novel, discrete, similar to those on the internet
Problems with idiosyncratic systems (custom setups)	low	novel, discrete, <i>not</i> similar to those the internet

Table 2: Conjectures about adoption by task

0.2 Additional issues.

1. **Why are LLMs used as advisors, not deputies?** It’s notable that LLMs are relatively rarely given autonomy to make a decision without human oversight. ChatGPT is mostly built as an advisor, it can’t take actions on your behalf, but it’s worth asking why that is. This could be accommodated by this model if most questions have large private components, such that ChatGPT’s answer is worse than the human’s, but the ChatGPT-augmented answer is better than the human’s.
2. **Relationship to time-savings.** This model quantifies the benefit from ChatGPT as the *accuracy* of an answer to a question. Much other literature on LLM measure the

value in time-savings, e.g. most RCTs of LLM augmentation (both in the laboratory and the field), and many self-reports. We could convert the accuracy increase to a time-savings if we say that people can spend more time to increase accuracy.

3. **How does ChatGPT differ from Google?** If we interpret ChatGPT’s training set as the content of the public internet then the same model could be applied to a search engine. We could distinguish ChatGPT in two ways: (1) the cost of consulting ChatGPT is significantly lower because it will give an answer immediately, instead of directing the human to another page where they have to skim the text; (2) ChatGPT can *interpolate* questions in its training data, e.g. it will give an answer to your question even if nobody has answered that question before, but the answer can be predicted from the answer to other questions.
4. **Dimensionality of the domain.** We might be able to extend the model to distinguish between two types of dimensionality: (1) surface dimensionality (how many letters it takes to express the question/prompt); (2) the latent dimensionality of the domain. In general the returns to experience will depend on the latent dimensionality of the domain: if the latent dimensionality is low then a few examples is enough to learn the patterns, if the latent dimensionality is high then error continuously decreases with experience (there’s a nice closed-form expression for this).
5. **High-dimensional answers.** Our model assumes *scalar* answers. In fact ChatGPT gives high-dimensional outputs. I think we can say some nice things here.
6. **Tacit knowledge.** For some ChatGPT prompts the user can *recognize* a correct answer, but cannot produce a correct answer themselves (AKA a generation-validation gap). E.g. asking for a picture, asking for a poem. The generation-validation asymmetry can be due either to (1) computational difficulty; (2) tacit knowledge.
7. **Routine questions.** For some ChatGPT prompts the user can clearly do the task themselves, without any extra information but it’s time-consuming: e.g. doing a simple mathematical operation, alphabetizing a list, typing out boilerplate code. These types of queries don’t fit our model so well.

1 Model

The State of the World and Questions. The state of the world is defined by a vector of p unobserved parameters, $\mathbf{w} \in \mathbb{R}^p$. A question is a vector of p binary features, $\mathbf{q} \in \{-1, 1\}^p$. The true answer to a question \mathbf{q} is a scalar a determined by the linear relationship:

$$a = \mathbf{q}'\mathbf{w} = \sum_{k=1}^p q_k w_k$$

Agents and Information. There is a set of agents, indexed by $i \in \mathcal{I}$. Each agent i possesses an information set \mathcal{D}_i , which consists of n_i questions they have previously encountered, along with their true answers. We can represent this information as a pair $(\mathbf{Q}_i, \mathbf{a}_i)$:

- \mathbf{Q}_i is an $n_i \times p$ matrix where each row is a question vector. Let the j -th question for agent i be $\mathbf{q}'_{i,j}$, so that:

$$\mathbf{Q}_i = \begin{bmatrix} \mathbf{q}'_{i,1} \\ \vdots \\ \mathbf{q}'_{i,n_i} \end{bmatrix} = \begin{bmatrix} q_{i,1,1} & \cdots & q_{i,1,p} \\ \vdots & \ddots & \vdots \\ q_{i,n_i,1} & \cdots & q_{i,n_i,p} \end{bmatrix}$$

- \mathbf{a}_i is an $n_i \times 1$ vector of the corresponding answers. The answers are generated according to the true model:

$$\mathbf{a}_i = \mathbf{Q}_i \mathbf{w}$$

Beliefs. All agents share a common prior belief about the state of the world, assuming the weights \mathbf{w} are drawn from a multivariate Gaussian distribution:

$$\mathbf{w} \sim N(\mathbf{0}, \Sigma)$$

where Σ is a $p \times p$ positive-semidefinite covariance matrix. A common assumption we will use is an isotropic prior, where $\Sigma = \sigma^2 \mathbf{I}_p$ for some scalar $\sigma^2 > 0$. This implies that, a priori, the weights are uncorrelated and have equal variance.

Given their information set \mathcal{D}_i , agent i forms a posterior belief about \mathbf{w} . When a new question \mathbf{q}_{new} arises, the agent uses their posterior distribution to form an estimate of the answer, $\hat{a}_{\text{new}} = \mathbf{q}'_{\text{new}} \mathbb{E}[\mathbf{w} \mid \mathcal{D}_i]$.

2 Propositions

Proposition 1 (Posterior over \mathbf{w} given \mathbf{Q} and \mathbf{a}). *The agent's posterior mean and variance will be:*

$$\begin{aligned} \hat{\mathbf{w}} &= \Sigma \mathbf{Q}^\top (\mathbf{Q} \Sigma \mathbf{Q}^\top)^{-1} \mathbf{a} \\ \Sigma_{|a} &= \Sigma - \Sigma \mathbf{Q}^\top (\mathbf{Q} \Sigma \mathbf{Q}^\top)^{-1} \mathbf{Q} \Sigma. \end{aligned}$$

Proof. The derivation follows from the standard formula for conditional Gaussian distributions. We begin by defining the joint distribution of the weights \mathbf{w} and the answers \mathbf{a} . The weights and answers are jointly Gaussian:

$$\begin{pmatrix} \mathbf{w} \\ \mathbf{a} \end{pmatrix} \sim N \left(\begin{pmatrix} \mathbf{0} \\ \mathbf{0} \end{pmatrix}, \begin{pmatrix} \Sigma & \Sigma \mathbf{Q}' \\ \mathbf{Q} \Sigma & \mathbf{Q} \Sigma \mathbf{Q}' \end{pmatrix} \right)$$

where the covariance terms are derived as follows:

- $Cov(\mathbf{w}, \mathbf{w}) = \Sigma$ (prior covariance)
- $Cov(\mathbf{a}, \mathbf{a}) = Cov(\mathbf{Q}\mathbf{w}, \mathbf{Q}\mathbf{w}) = \mathbf{Q}Cov(\mathbf{w}, \mathbf{w})\mathbf{Q}' = \mathbf{Q}\Sigma\mathbf{Q}'$
- $Cov(\mathbf{w}, \mathbf{a}) = Cov(\mathbf{w}, \mathbf{Q}\mathbf{w}) = Cov(\mathbf{w}, \mathbf{w})\mathbf{Q}' = \Sigma\mathbf{Q}'$

The conditional mean $E[\mathbf{w}|\mathbf{a}]$ is given by the formula:

$$E[\mathbf{w}|\mathbf{a}] = E[\mathbf{w}] + Cov(\mathbf{w}, \mathbf{a})Var(\mathbf{a})^{-1}(\mathbf{a} - E[\mathbf{a}])$$

Substituting the values from our model ($E[\mathbf{w}] = \mathbf{0}$, $E[\mathbf{a}] = \mathbf{0}$):

$$\hat{\mathbf{w}} = \mathbf{0} + (\Sigma\mathbf{Q}')(\mathbf{Q}\Sigma\mathbf{Q}')^{-1}(\mathbf{a} - \mathbf{0}) = \Sigma\mathbf{Q}'(\mathbf{Q}\Sigma\mathbf{Q}')^{-1}\mathbf{a}$$

This gives us the posterior mean of the weights. The posterior covariance is given by:

$$Var(\mathbf{w}|\mathbf{a}) = Var(\mathbf{w}) - Cov(\mathbf{w}, \mathbf{a})Var(\mathbf{a})^{-1}Cov(\mathbf{a}, \mathbf{w}) = \Sigma - \Sigma\mathbf{Q}'(\mathbf{Q}\Sigma\mathbf{Q}')^{-1}\mathbf{Q}\Sigma.$$

□

Proposition 2 (Expected error for a given question). *The expected squared error for a new question \mathbf{q} is:*

$$\mathbb{E}[(\mathbf{q}'(\mathbf{w} - \hat{\mathbf{w}}))^2] = \mathbf{q}'\Sigma_{|\mathbf{a}}\mathbf{q}$$

For an isotropic prior where $\Sigma = \sigma^2\mathbf{I}$, the error is proportional to the squared distance of \mathbf{q} from the subspace spanned by the previously seen questions in \mathbf{Q} :

$$\mathbb{E}[(\mathbf{q}'(\mathbf{w} - \hat{\mathbf{w}}))^2] = \sigma^2\|(\mathbf{I} - \mathbf{P}_{\mathbf{Q}})\mathbf{q}\|^2$$

where $\mathbf{P}_{\mathbf{Q}}$ is the projection matrix onto the row-span of \mathbf{Q} .

Proof. The prediction error is $\mathbf{q}'\mathbf{w} - \mathbf{q}'\hat{\mathbf{w}} = \mathbf{q}'(\mathbf{w} - \hat{\mathbf{w}})$. The expected squared error is the variance of this prediction error.

$$\begin{aligned}\mathbb{E}[(\mathbf{q}'(\mathbf{w} - \hat{\mathbf{w}}))^2] &= \mathbb{E}[\mathbf{q}'(\mathbf{w} - \hat{\mathbf{w}})(\mathbf{w} - \hat{\mathbf{w}})'\mathbf{q}] \\ &= \mathbf{q}'\mathbb{E}[(\mathbf{w} - \hat{\mathbf{w}})(\mathbf{w} - \hat{\mathbf{w}})']\mathbf{q} \\ &= \mathbf{q}'\text{Var}(\mathbf{w} \mid \mathbf{a})\mathbf{q} = \mathbf{q}'\Sigma_{|\mathbf{a}}\mathbf{q}\end{aligned}$$

This proves the first part of the proposition. For the second part, we assume an isotropic prior $\Sigma = \sigma^2\mathbf{I}$. Substituting this into the expression for $\Sigma_{|\mathbf{a}}$ from Proposition 1:

$$\begin{aligned}\Sigma_{|\mathbf{a}} &= \sigma^2\mathbf{I} - (\sigma^2\mathbf{I})\mathbf{Q}'(\mathbf{Q}(\sigma^2\mathbf{I})\mathbf{Q}')^{-1}\mathbf{Q}(\sigma^2\mathbf{I}) \\ &= \sigma^2\mathbf{I} - \sigma^4\mathbf{Q}'(\sigma^2\mathbf{Q}\mathbf{Q}')^{-1}\mathbf{Q} \\ &= \sigma^2\mathbf{I} - \sigma^4(\sigma^2)^{-1}\mathbf{Q}'(\mathbf{Q}\mathbf{Q}')^{-1}\mathbf{Q} \\ &= \sigma^2(\mathbf{I} - \mathbf{Q}'(\mathbf{Q}\mathbf{Q}')^{-1}\mathbf{Q})\end{aligned}$$

Let $\mathbf{P}_{\mathbf{Q}} = \mathbf{Q}'(\mathbf{Q}\mathbf{Q}')^{-1}\mathbf{Q}$, which is the projection matrix onto the row space of \mathbf{Q} . Then $\Sigma_{|\mathbf{a}} = \sigma^2(\mathbf{I} - \mathbf{P}_{\mathbf{Q}})$. The expected squared error is:

$$\mathbb{E}[(\mathbf{q}'(\mathbf{w} - \hat{\mathbf{w}}))^2] = \mathbf{q}'\sigma^2(\mathbf{I} - \mathbf{P}_{\mathbf{Q}})\mathbf{q} = \sigma^2\mathbf{q}'(\mathbf{I} - \mathbf{P}_{\mathbf{Q}})\mathbf{q}$$

Since $\mathbf{I} - \mathbf{P}_{\mathbf{Q}}$ is an idempotent projection matrix, $\mathbf{q}'(\mathbf{I} - \mathbf{P}_{\mathbf{Q}})\mathbf{q} = \mathbf{q}'(\mathbf{I} - \mathbf{P}_{\mathbf{Q}})'(\mathbf{I} - \mathbf{P}_{\mathbf{Q}})\mathbf{q} = \|(\mathbf{I} - \mathbf{P}_{\mathbf{Q}})\mathbf{q}\|^2$. Thus,

$$\mathbb{E}[(\mathbf{q}'(\mathbf{w} - \hat{\mathbf{w}}))^2] = \sigma^2\|(\mathbf{I} - \mathbf{P}_{\mathbf{Q}})\mathbf{q}\|^2$$

□

Proposition 3 (Error decreases with more independent questions). *The average expected squared error over all possible new questions \mathbf{q} decreases linearly with the number of linearly independent questions in the training set \mathbf{Q} . Specifically, with an isotropic prior $\Sigma = \sigma^2\mathbf{I}$, the average error is:*

$$\mathbb{E}_{\mathbf{q}}[\text{error}(\mathbf{q})] = \sigma^2(p - \text{rank}(\mathbf{Q}))$$

where the expectation is taken over new questions \mathbf{q} with i.i.d. components drawn uniformly from $\{-1, 1\}$.

Proof. The proof proceeds in two steps. First, we write the expression for the error for a given new question \mathbf{q} . Second, we average this error over the distribution of all possible questions.

1. **Predictive error for a fixed \mathbf{q} .** From Proposition 2, the expected squared error for a specific new question \mathbf{q} , given an isotropic prior $\Sigma = \sigma^2 \mathbf{I}$, is:

$$\text{error}(\mathbf{q}) = \mathbb{E}[(\mathbf{q}'(\mathbf{w} - \hat{\mathbf{w}}))^2] = \sigma^2 \mathbf{q}'(\mathbf{I} - \mathbf{P}_Q)\mathbf{q}$$

where $\mathbf{P}_Q = \mathbf{Q}'(\mathbf{Q}\mathbf{Q}')^{-1}\mathbf{Q}$ is the projection matrix onto the row-span of \mathbf{Q} .

2. **Average over random new questions.** We now take the expectation of this error over the distribution of new questions \mathbf{q} . The components of \mathbf{q} are i.i.d. uniform on $\{-1, 1\}$, which implies that $\mathbb{E}[\mathbf{q}] = \mathbf{0}$ and $\mathbb{E}[\mathbf{q}\mathbf{q}'] = \mathbf{I}_p$. The average error is:

$$\begin{aligned} \mathbb{E}_q[\text{error}(\mathbf{q})] &= \mathbb{E}_q[\sigma^2 \mathbf{q}'(\mathbf{I} - \mathbf{P}_Q)\mathbf{q}] \\ &= \sigma^2 \mathbb{E}_q[\text{tr}(\mathbf{q}'(\mathbf{I} - \mathbf{P}_Q)\mathbf{q})] \\ &= \sigma^2 \mathbb{E}_q[\text{tr}((\mathbf{I} - \mathbf{P}_Q)\mathbf{q}\mathbf{q}')] \\ &= \sigma^2 \text{tr}((\mathbf{I} - \mathbf{P}_Q)\mathbb{E}_q[\mathbf{q}\mathbf{q}']) \\ &= \sigma^2 \text{tr}(\mathbf{I} - \mathbf{P}_Q) \\ &= \sigma^2 (\text{tr}(\mathbf{I}) - \text{tr}(\mathbf{P}_Q)) \end{aligned}$$

The trace of the identity matrix is p . The trace of a projection matrix is the dimension of the subspace it projects onto, so $\text{tr}(\mathbf{P}_Q) = \text{rank}(\mathbf{Q})$. Thus, the average error is:

$$\mathbb{E}_q[\text{error}(\mathbf{q})] = \sigma^2 (p - \text{rank}(\mathbf{Q}))$$

Since the rank of \mathbf{Q} increases with each linearly independent question added, the average error decreases linearly until $\text{rank}(\mathbf{Q}) = p$, at which point it becomes zero.

□

Proposition 4 (Two-stage updating with agents 1 and 2). *Consider two agents who share an isotropic prior $\mathbf{w} \sim N(\mathbf{0}, \sigma^2 \mathbf{I}_p)$.*

- Agent 1 observes data $(\mathbf{Q}_1, \mathbf{a}_1)$ and forms the posterior mean

$$\hat{\mathbf{w}}_1 = \mathbf{Q}_1^\top (\mathbf{Q}_1 \mathbf{Q}_1^\top)^{-1} \mathbf{a}_1, \quad \mathbf{P}_1 := \mathbf{Q}_1^\top (\mathbf{Q}_1 \mathbf{Q}_1^\top)^{-1} \mathbf{Q}_1.$$

- Agent 2 observes data $(\mathbf{Q}_2, \mathbf{a}_2)$ and forms the posterior mean

$$\hat{\mathbf{w}}_2 = \mathbf{Q}_2^\top (\mathbf{Q}_2 \mathbf{Q}_2^\top)^{-1} \mathbf{a}_2, \quad \mathbf{P}_2 := \mathbf{Q}_2^\top (\mathbf{Q}_2 \mathbf{Q}_2^\top)^{-1} \mathbf{Q}_2.$$

- A new question $\mathbf{q} \in \{-1, 1\}^p$ arrives. Agent 1 announces the estimate $\hat{a}_1 = \mathbf{q}'\hat{\mathbf{w}}_1$. Let

$$\mu_2 = \mathbf{q}'\hat{\mathbf{w}}_2, \quad \sigma_2^2 = \sigma^2 \mathbf{q}'(\mathbf{I} - \mathbf{P}_2)\mathbf{q}, \quad (1)$$

$$\mu_{2,1} = \mathbf{q}'\mathbf{P}_1\hat{\mathbf{w}}_2, \quad \sigma_{21} = \sigma^2 \mathbf{q}'(\mathbf{I} - \mathbf{P}_2)\mathbf{P}_1^\top \mathbf{q}, \quad (2)$$

$$\sigma_{1|2}^2 = \sigma^2 \mathbf{q}'\mathbf{P}_1(\mathbf{I} - \mathbf{P}_2)\mathbf{P}_1^\top \mathbf{q}, \quad \kappa = \frac{\sigma_{21}}{\sigma_{1|2}^2}. \quad (3)$$

Then, the posterior distribution of the true answer $a = \mathbf{q}'\mathbf{w}$ for agent 2 before seeing \hat{a}_1 is $N(\mu_2, \sigma_2^2)$, and after observing \hat{a}_1 it is

$$a \mid \hat{a}_1, \mathbf{a}_2 \sim N(\mu_{2|1}, \sigma_{2|1}^2), \quad \mu_{2|1} = \mu_2 + \kappa(\hat{a}_1 - \mu_{2,1}), \quad \sigma_{2|1}^2 = \sigma_2^2 - \kappa \sigma_{21}.$$

Proof. The estimate of agent 1 is a linear function of the true weights \mathbf{w} , since $\mathbf{a}_1 = \mathbf{Q}_1\mathbf{w}$, so $\hat{a}_1 = \mathbf{q}'\hat{\mathbf{w}}_1 = \mathbf{q}'\mathbf{Q}_1^\top(\mathbf{Q}_1\mathbf{Q}_1^\top)^{-1}\mathbf{a}_1 = \mathbf{q}'\mathbf{P}_1\mathbf{w}$.

Conditioning on agent 2's data, the posterior for \mathbf{w} is $N(\hat{\mathbf{w}}_2, \Sigma_2)$ with $\Sigma_2 = \sigma^2(\mathbf{I} - \mathbf{P}_2)$. The pair (a, \hat{a}_1) is therefore jointly Gaussian, since both are linear functions of \mathbf{w} . Their joint distribution conditional on agent 2's data has:

- $E[a|\mathbf{a}_2] = \mathbf{q}'\hat{\mathbf{w}}_2 = \mu_2$
- $E[\hat{a}_1|\mathbf{a}_2] = \mathbf{q}'\mathbf{P}_1\hat{\mathbf{w}}_2 = \mu_{2,1}$
- $\text{Var}(a|\mathbf{a}_2) = \mathbf{q}'\sigma^2(\mathbf{I} - \mathbf{P}_2)\mathbf{q} = \sigma_2^2$
- $\text{Var}(\hat{a}_1|\mathbf{a}_2) = \mathbf{q}'\mathbf{P}_1\sigma^2(\mathbf{I} - \mathbf{P}_2)\mathbf{P}_1^\top\mathbf{q} = \sigma_{1|2}^2$
- $\text{Cov}(a, \hat{a}_1|\mathbf{a}_2) = \mathbf{q}'\sigma^2(\mathbf{I} - \mathbf{P}_2)\mathbf{P}_1^\top\mathbf{q} = \sigma_{21}$

So the covariance matrix of (a, \hat{a}_1) conditional on agent 2's data is:

$$\begin{pmatrix} \sigma_2^2 & \sigma_{21} \\ \sigma_{21} & \sigma_{1|2}^2 \end{pmatrix}$$

For any joint Gaussian vector, the conditional distribution of the first component given the second is again Gaussian with

$$\mu_{2|1} = \mu_2 + \frac{\sigma_{21}}{\sigma_{1|2}^2}(\hat{a}_1 - \mu_{2,1}), \quad \sigma_{2|1}^2 = \sigma_2^2 - \frac{\sigma_{21}^2}{\sigma_{1|2}^2}.$$

Identifying $\kappa = \sigma_{21}/\sigma_{1|2}^2$ yields the stated result. □

Proposition 5 (Conditions for valuable two-stage updating). *In the setting of Proposition 4, consulting agent 1 provides value to agent 2 if and only if:*

$$\mathbf{q}'(\mathbf{I} - \mathbf{P}_2)\mathbf{P}_1^\top\mathbf{q} \neq 0$$

When this condition holds:

- The posterior mean changes: $\mu_{2|1} \neq \mu_2$
- The posterior variance decreases: $\sigma_{2|1}^2 < \sigma_2^2$

When this condition fails, consulting agent 1 provides no additional information: $\mu_{2|1} = \mu_2$ and $\sigma_{2|1}^2 = \sigma_2^2$.

Proof. From Proposition 4, the change in the posterior mean is $\mu_{2|1} - \mu_2 = \kappa(\hat{a}_1 - \mu_{2,1})$, and the change in posterior variance is $\sigma_2^2 - \sigma_{2|1}^2 = \kappa\sigma_{21}$, where $\kappa = \frac{\sigma_{21}}{\sigma_{1|2}^2}$ and $\sigma_{21} = \sigma^2 \mathbf{q}'(\mathbf{I} - \mathbf{P}_2)\mathbf{P}_1^\top \mathbf{q}$.
 If $\sigma_{21} = 0$, then $\kappa = 0$, so $\mu_{2|1} = \mu_2$ and $\sigma_{2|1}^2 = \sigma_2^2$.
 If $\sigma_{21} \neq 0$, then $\kappa \neq 0$ (since $\sigma_{1|2}^2 \geq 0$ with equality only when $\mathbf{P}_1(\mathbf{I} - \mathbf{P}_2) = \mathbf{0}$, which implies $\sigma_{21} = 0$). In this case, both the mean and variance will generally change unless $\hat{a}_1 = \mu_{2,1}$, which occurs with probability zero.
 Therefore, two-stage updating provides value if and only if $\sigma_{21} = \sigma^2 \mathbf{q}'(\mathbf{I} - \mathbf{P}_2)\mathbf{P}_1^\top \mathbf{q} \neq 0$. \square

The intuition behind Proposition 5 is straightforward: **it is worthwhile to consult another agent if and only if the component of the question that you don't understand overlaps with the other agent's area of expertise.**

More precisely:

- $(\mathbf{I} - \mathbf{P}_2)\mathbf{q}$ represents the *residual* of the question after projecting it onto agent 2's own experience. This is the part of the question that agent 2 finds novel or unfamiliar.
- $\mathbf{P}_1^\top \mathbf{q}$ represents the component of the question that lies within agent 1's area of expertise (the row space of their experience matrix \mathbf{Q}_1).
- The condition $\mathbf{q}'(\mathbf{I} - \mathbf{P}_2)\mathbf{P}_1^\top \mathbf{q} \neq 0$ requires that these two components are not orthogonal—there must be some overlap between what agent 2 doesn't know and what agent 1 does know.

This formalizes the intuitive notion that collaboration is valuable when agents have *complementary* rather than identical or completely unrelated knowledge. If agent 1's expertise is orthogonal to the unfamiliar aspects of the question for agent 2, then agent 1's opinion provides no useful information. Conversely, if there is overlap between agent 2's knowledge gaps and agent 1's strengths, then consultation becomes valuable.

In the context of the ChatGPT model, this suggests that an AI assistant is most valuable for questions where:

1. The question contains elements that are novel to the human user (large $\|(\mathbf{I} - \mathbf{P}_2)\mathbf{q}\|$)
2. These novel elements fall within the AI's training domain (non-zero projection onto the AI's knowledge space)

3 References
